

Introduction to Cluster Computing

UVA – AMSTERDAM – 01.11.2019

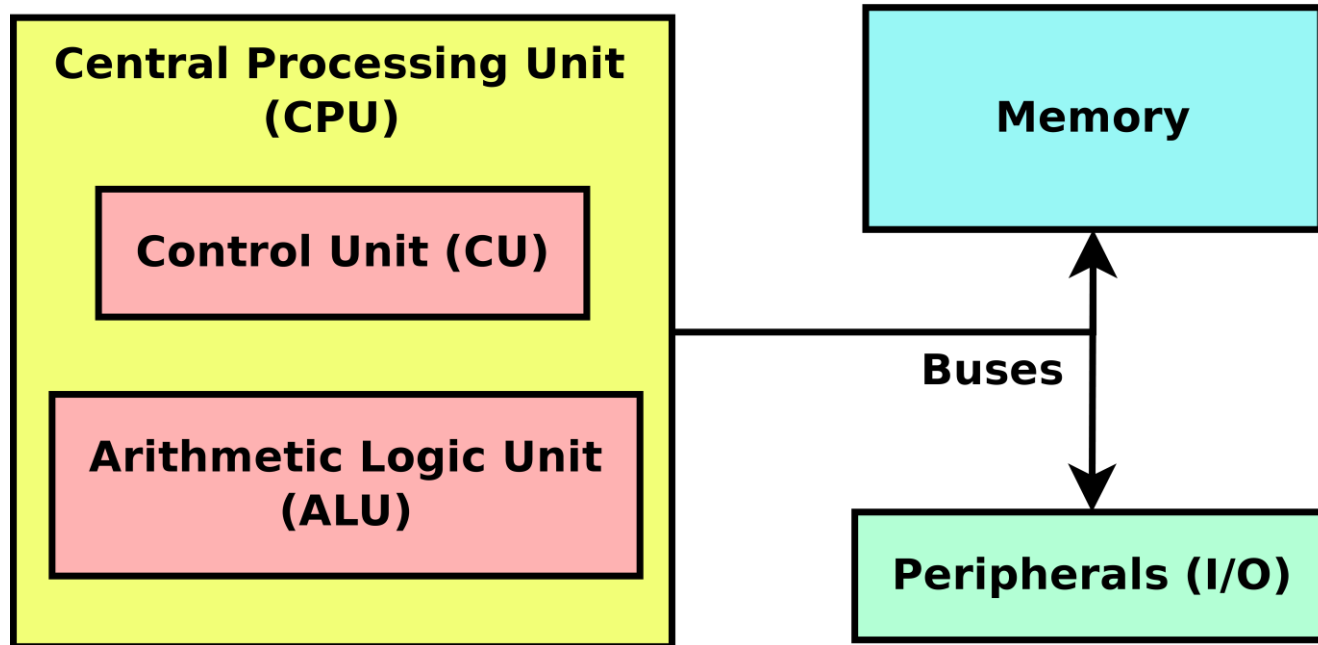


Carlos Teijeiro Barjas


Advisor Cluster Computing carlos.teijeiro@surfsara.nl




A computer is...




A larger computer usually is...




A larger computer usually is...



A larger computer usually is...



A larger computer usually is...



Outline

- **SURFsara facilities**
- **Running jobs with SLURM**
- **Hands-on exercises on the Lisa GPU island**

Outline

- **SURFsara facilities**
- **Running jobs with SLURM**
- **Hands-on exercises on the Lisa GPU island**

What do we do at SURFsara?


- **Regular user support for the supercomputers Cartesius and Lisa**
 - Typical effort: from a few minutes to a couple of days
- **Application enabling for Dutch Compute Challenge Projects**
 - Potential effort by SURFsara staff: 1 to 6 person months per project
- **Performance improvement of applications**
 - Typically meant for promising user applications
 - Potential effort by SURFsara staff: 3 to 6 person months per project
- **Support for PRACE applications**
 - PRACE offers access to European systems
 - SURFsara participates in PRACE support in application enabling
- **Visualization projects**
- **Training and workshops on demand**
- **Please contact SURFsara at helpdesk@surfsara.nl**

Historical performance of Dutch supercomputers

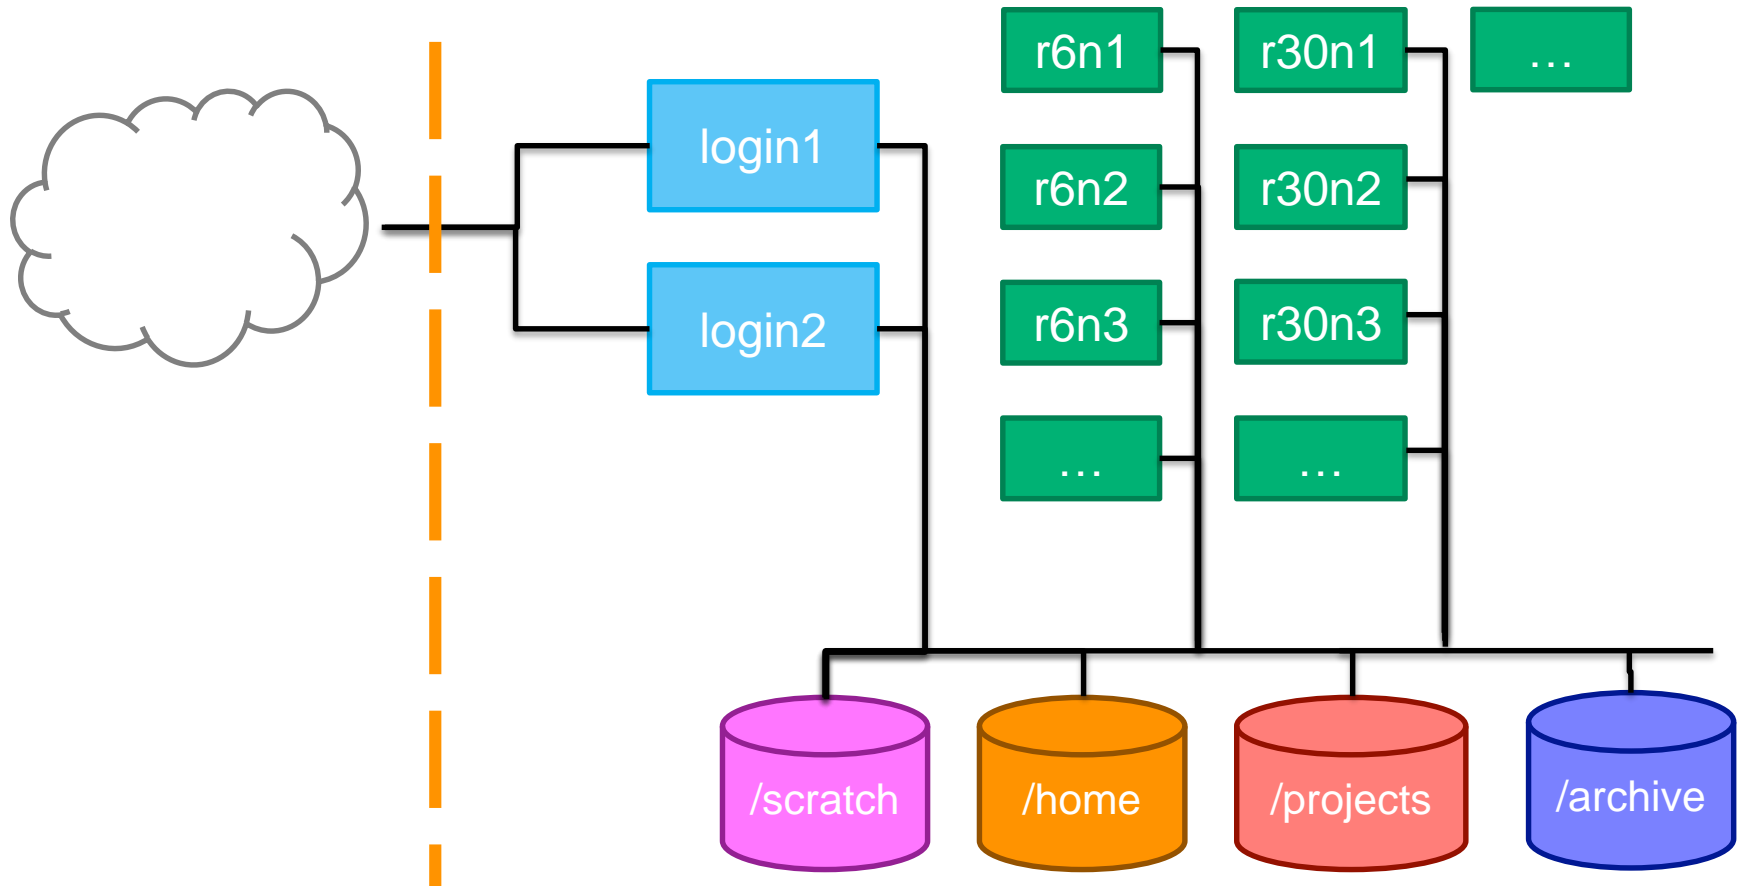
| Year | Machine | R_{peak} GFlop/s | kW | GFlop/s / kW |
|------|---------------------------------|------------------------------|-------|-----------------|
| 1984 | CDC Cyber 205 1-pipe | 0.1 | 250 | 0.0004 |
| 1988 | CDC Cyber 205 2-pipe | 0.2 | 250 | 0.0008 |
| 1991 | Cray Y-MP/4128 | 1.33 | 200 | 0.0067 |
| 1994 | Cray C98/4256 | 4 | 300 | 0.0133 |
| 1997 | Cray C916/121024 | 12 | 500 | 0.024 |
| 2000 | SGI Origin 3800 | 1,024 | 300 | 3.4 |
| 2004 | SGI Origin 3800 +SGI Altix 3700 | 3,200 | 500 | 6.4 |
| 2007 | IBM p575 Power5+ | 14,592 | 375 | 40 |
| 2008 | IBM p575 Power6 | 62,566 | 540 | 116 |
| 2009 | IBM p575 Power6 | 64,973 | 560 | 116 |
| 2013 | Bull bullx DLC | 250,000 | 260 | 962 |
| 2014 | Bull bullx DLC | >1,000,000 | >520 | 1923 |
| 2017 | Bull bullx DLC + KNL | > 1,800,000 | | |
| 2016 | Raspberry PI 3 (35 euro) | 0.44 | 0.004 | 110 |



Schematic overview of Cartesius & Lisa



Specific example: Lisa architecture



Lisa – Nodes

- The two login nodes are of type E5-2650 v2

| Number | Processor Type | Clock | Scratch | Memory | Cache | Cores | GPUs | Interconnect |
|--------|-----------------------------|----------|----------------|-------------------------|----------|-------|--|-----------------------|
| 32 | E5-2650 v2 | 2.60 GHz | 870 GB | 32 GB QPI 8.00 GT/s | 20 MB | 16 | - | 10 Gbit/s ethernet |
| 280 | E5-2650 v2 | 2.60 GHz | 870 GB | 64 GB QPI 8.00 GT/s | 20 MB | 16 | - | 10 Gbit/s ethernet |
| 32 | E5-2650 v2 | 2.60 GHz | 870 GB | 64 GB QPI 8.00 GT/s | 20 MB | 16 | - | Mellanox FDR |
| 23 | Bronze 3104 | 1.70 GHz | 1.5 TB NVME | 256 GB UPI 10.4 GT/s | 8.25 MB | 12 | 4 x GeForce 1080Ti, 11GB GDDR5X | 40 Gbit/s ethernet |
| 96 | Silver 4110 | 2.10 GHz | 1.8 TB | 64 GB UPI 9.6 GT/s | 11 MB | 16 | - | 10 Gbit/s ethernet |
| 1 | E7-8857 v2 | 3.00 GHz | 13 TB | 1 TB QPI 8.00 GT/s | 30 MB | 48 | - | 10 Gbit/s ethernet |
| 1 | Gold 6126 | 2.60 GHz | 11 TB | 2 TB UPI 10.4 GT/s | 19.25 MB | 48 | - | 40 Gbit/s ethernet |

Lisa – Nodes

CPU nodes

Total number of cores 7484

Total amount of memory 26 TB

Total peak performance 149 TFlop/sec

Disk space 400 TB for the home file systems

Mellanox InfiniBand network - FDR: 56 Gbit/sec Latency FDR: 1.3 μ sec

GPU nodes

Total number of cores 276

Total amount of memory 5.9 TB

Total peak performance (SP) 1.026 TFlop/sec

Total peak performance (DP) 32 TFlop/sec

Disk space 400 TB for the home file systems

Interconnect 40 Gbit/s ethernet

Cartesius & Lisa – File systems

- **/home/user**
- User home directory (quota - currently 200GB)
- Backed up
- Meant for storage of important files (sources, scripts, input and output data)
- Not the fastest file system
- **/scratch**
- Cartesius: /scratch-local & /scratch-shared (quota – currently 8 TB)
- Lisa: /scratch (quota – depends on disk size)
- Not backed up
- Meant for temporary storage (during running of a job and shortly thereafter)
- The fastest file systems on Cartesius & Lisa

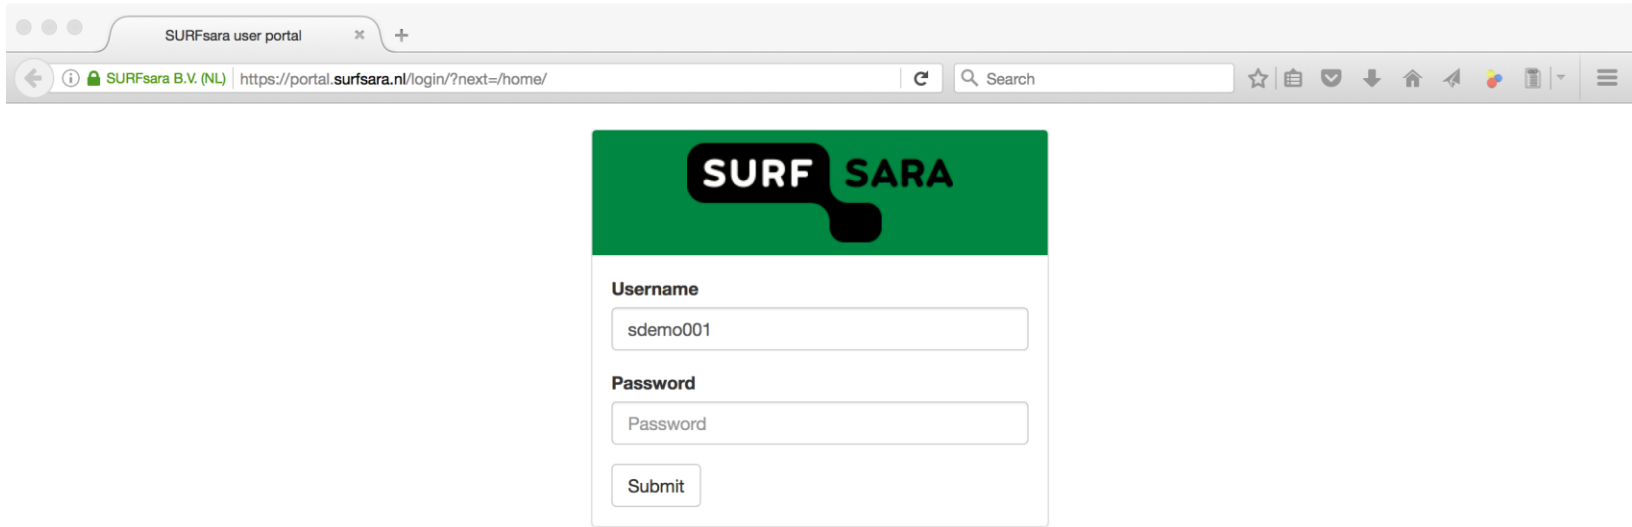
Cartesius & Lisa – File systems

- **/archive**
 - Connected to the tape robot (quota – virtually unlimited)
 - Backed up
 - Meant for long term storage of files, zipped, tarred, combined into small number of files
 - Slow – especially when retrieving “old” data
 - Not available to worker nodes
- **/project**
 - On Cartesius: large and fast, On Lisa: large, but not so fast
 - For special projects requiring lots of space (quota – as much as needed/possible)
 - Not backed up
 - Meant for special projects
 - On Cartesius: comparable in speed with /scratch. On Lisa: comparable to /home.

Outline

- **SURFsara facilities**
- **Running jobs with SLURM**
- **Hands-on exercises on the Lisa GPU island**

Running jobs: first change your password



The screenshot shows a web browser window with the title "SURFsara user portal". The address bar displays "https://portal.surfsara.nl/login/?next=/home/" with a search icon and a "Search" input field. The page content features a green header with the "SURF SARA" logo. Below the header is a login form with the following fields and buttons:

- Username**: A text input field containing "sdemo001".
- Password**: A text input field containing "Password".
- Submit**: A button to submit the login information.

Running jobs: now connect to Lisa

- Connect to Lisa from MobaXterm or terminal

```
user@localmachine:~$ ssh lgpu0XXX@lisa.surfsara.nl  
Password:  
lgpu0XXX@login1:~$
```

- If you are not familiar with Linux commands, try to follow some tutorial. For example:

<https://swcarpentry.github.io/shell-novice/>

Running jobs: how-to

Remember: when you log in, you log in a *login node*.

The SLURM scheduler will distribute work to *batch nodes*

Workflow:

1. **You** upload your data from your computer to the cluster system
2. **You** create a job script with the work steps
3. **You** submit the job script to the scheduler
4. **The scheduler** looks for available computers to run your work
5. When a batch node with the requirements you specified becomes available, your work runs
6. When the job finishes, you can get an e-mail to inform you
7. When the job is finished, **you** download the results to your computer

Running jobs: first example

```
#!/bin/bash
#SBATCH -t 0:02:00
#SBATCH -n 1
#SBATCH -p gpu_shared_course

echo "Let's do some work"
sleep 20
echo "All work done."
```

- Create a text file with *exactly* the first lines; name the file “job.sh”
- Submit this job with “**sbatch** job.sh”
- Use “**squeue -u username**” to look at the status
- Use “**scontrol show job *jobid***” to find out when your job will run
- Look at your home-directory to see what happens there; look at the files.
- Which files were created? Look at those files.

Running jobs: useful commands

- | | |
|--|-------------------------------------|
| sbatch <jobscript> | - submit a job to the scheduler |
| squeue <jobid> | - inspect the status of job <jobid> |
| squeue -u <userid> | - inspect all jobs of user <userid> |
| scancel <jobid> | - cancel job <jobid> before it runs |
| scontrol show job <jobid> | - show estimated job start |

Running jobs: best practices

- **Give the scheduler a realistic *walltime* estimate**
- **Your home directory is slow. Use \$TMPDIR: it is a temporary directory in /scratch created for your job (but then copy the results back to your /home !!!)**
- **Load software modules as part of your job script – this improves reproducibility**
- **Run parallel versions of your programs**

Anatomy of a job script

Job scripts consist of:

- the “shebang” line: `#!/bin/bash`
- scheduler directives: `#SBATCH ...`
- loading software modules: `module load ...`
- setting environment: `export VAR=...`
- preparing input
- running your program
- saving output

Example: a real job script

```
#!/bin/bash
#SBATCH -t 0:20:00
#SBATCH -N 1 -c 24

module load python/3.5.2

cp -r $HOME/run3 $TMPDIR

cd $TMPDIR/run3
python myscript.py input.dat

mkdir -p $HOME/run3/results
cp result.dat run3.log $HOME/run3/results
```

Module management: useful commands

- | | |
|----------------------------------|---------------------------------------|
| module avail | - available modules in the system |
| module load <mod> | - load <mod> in the shell environment |
| module list | - show a list of all loaded modules |
| module unload <mod> | - remove <mod> from the environment |
| module purge | - unload all modules |
| module whatis <mod> | - show information about <mod> |

Everything about jobs: user info pages

Go to:

<https://userinfo.surfsara.nl>

Click on the corresponding system:

- **Cartesius: Usage → Batch Usage (jobs)**
- **Lisa: User guide → 4. Creating and running jobs**

Outline

- **SURFsara facilities**
- **Running jobs with SLURM**
- **Hands-on exercises on the Lisa GPU island**

Hands-on: downloading the dataset

- We are going to use the CIFAR10 dataset for machine learning tests on Lisa GPUs
 - Contains 50000 training and 10000 test images for classification
 - Web page: <https://www.cs.toronto.edu/~kriz/cifar.html>
 - Use **wget** to download the dataset directly to your /home directory on Lisa using the login node

```
lgpu0XXX@login1:~$ wget \  
https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
```

Hands-on: installing PyTorch & TensorFlow

- Two local (user) installations are necessary
 - Load the required modules for Python and CUDA functions
 - Export the library path to load additional libraries
 - Use pip3 to perform the local installations

```
lgpu0XXX@login1:~$ module load Python/3.6.3-foss-2017b
lgpu0XXX@login1:~$ module load cuDNN/7.0.5-CUDA-9.0.176
lgpu0XXX@login1:~$ module load NCCL/2.0.5-CUDA-9.0.176
lgpu0XXX@login1:~$ export LD_LIBRARY_PATH=/hpc/eb/Debian9/
cuDNN/7.1-CUDA-8.0.44-GCCcore-5.4.0/lib64:$LD_LIBRARY_PATH
lgpu0XXX@login1:~$ pip3 install --user torch torchvision
lgpu0XXX@login1:~$ pip3 install --user tensorflow_gpu
```

Hands-on: connecting to login-gpu on Lisa

- Now we connect to the login-gpu node of Lisa !!!
 - NOTE: the previous step can be done on this login-gpu node too, but the main login nodes of Lisa are faster for **wget** and local installations with **pip**

```
lgpu0XXX@login1:~$ logout
user@localmachine:~$ ssh lgpu0XXX@login-gpu.lisa.surfsara.nl
Password:
lgpu0XXX@login-gpu1:~$
```

Example hands-on: PyTorch

- Download the examples directory for PyTorch
- Copy the CIFAR10 dataset to the dcgan folder to run a generative adversarial network (GAN)
- Use that folder as working directory and edit your batch script there

```
lgpu0XXX@login-gpu1:~$ git clone \  
https://github.com/pytorch/examples.git pytorch_examples  
lgpu0XXX@login-gpu1:~$ cp cifar-10-python.tar.gz \  
pytorch_examples/dcgan/  
lgpu0XXX@login-gpu1:~$ cd pytorch_examples/dcgan/  
lgpu0XXX@login-gpu1:~$ nano pytorch.job
```


Example hands-on: PyTorch

```
#!/bin/bash
#SBATCH --job-name=example
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=3
#SBATCH --ntasks-per-node=1
#SBATCH --time=1:00:00
#SBATCH --mem=60000M
#SBATCH --partition=gpu_shared_course
#SBATCH --gres=gpu:1

module load pre2019
module load Python/3.6.3-foss-2017b
module load cuDNN/7.0.5-CUDA-9.0.176
module load NCCL/2.0.5-CUDA-9.0.176
export LD_LIBRARY_PATH=/hpc/eb/Debian9/cuDNN/7.1-CUDA-8.0.44-
GCCcore-5.4.0/lib64:$LD_LIBRARY_PATH

srun python3 -u main.py --dataset=cifar10 --dataroot=. --cuda
```

Example hands-on: PyTorch

- Submit the job...

```
lgpu0XXX@login-gpu1:~$ sbatch pytorch.job
```

- ... and now you can also try the TensorFlow example on your own!
- (the guidelines are in your home folders)

Introduction to Cluster Computing

UVA – AMSTERDAM – 01.11.2019



Carlos Teijeiro Barjas

Advisor Cluster Computing carlos.teijeiro@surfsara.nl

